

In this project, we will train a GNN to perform link prediction on a heterogenous graph from the Spotify Million Playlists dataset.

## Import libraries

```
!pip install torch numpy matplotlib torch_geometric torch-scatter torcheval
# https://stackoverflow.com/a/73534928
import torch
!pip install torch-sparse -f https://data.pyg.org/whl/torch-{torch.__version__}.html

import tqdm
import torch
import torch_geometric
import numpy as np
import time
import matplotlib.pyplot as plt
import torch_geometric.transforms as T
from torcheval.metrics import BinaryAccuracy

import itertools
import time
import networkx as nx
import json
import os

import networkx as nx
from tqdm import tqdm
```

## Configuration

Dataset files

```
from google.colab import drive
drive.mount('/content/drive')

dataset_path = "drive/MyDrive/spotify_million_playlist_dataset/data"

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

## Load graph

```
def load_graph(dataset_path=dataset_path):
    """Load a nx.Graph from disk."""
    filenames = os.listdir(dataset_path)
    G = nx.DiGraph()
    for i in tqdm(range(len(filenames)), unit="files"):
        with open(os.path.join(dataset_path, filenames[i])) as json_file:
            playlists = json.load(json_file)["playlists"]
            for playlist in playlists:
                playlist_name = f"spotify:playlist:{playlist['pid']}"
                G.add_node(
                    playlist_name,
                    node_type="playlist",
                    num_followers=playlist["num_followers"],
                    num_tracks=playlist["num_tracks"],
                    num_artists=playlist["num_artists"],
                    num_albums=playlist["num_albums"],
                    duration_ms=playlist["duration_ms"],
                    collaborative=playlist["collaborative"],
                    num_edits=playlist["num_edits"]
                )
            for track in playlist["tracks"]:
                G.add_node(track["track_uri"], node_type="track", duration=track["duration_ms"])
                G.add_node(track["album_uri"], node_type="album")
                G.add_node(track["artist_uri"], node_type="artist")

                G.add_edge(track["track_uri"], playlist_name, edge_type="track-playlist")
                G.add_edge(track["track_uri"], track["album_uri"], edge_type="track-album")
                G.add_edge(track["track_uri"], track["artist_uri"], edge_type="track-artist")

    return G

def nx2hetero(G):
```

```

"""Convert a nx.Graph into a torch_geometric.data.HeteroData object."""
ids_by_type = {
    "playlist": {},
    "track": {},
    "artist": {},
    "album": {}
}

def node_id(node_type, id):
    d = ids_by_type[node_type]
    if id not in d:
        d[id] = len(d)
    return d[id]

node_features_by_type = {
    "playlist": [],
    "track": [],
    "artist": [],
    "album": []
}

# comment:
# {
#     "name": "musical",
#     "collaborative": "false",
#     "pid": 5,
#     "modified_at": 1493424000,
#     "num_albums": 7,
#     "num_tracks": 12,
#     "num_followers": 1,
#     "num_edits": 2,
#     "duration_ms": 2657366,
#     "num_artists": 6,
#     "tracks": [
#         {
#             "pos": 0,
#             "artist_name": "Degiheugi",
#             "track_uri": "spotify:track:7vqa3sDmtEaVJ2gcvxtRID",
#             "artist_uri": "spotify:artist:3V2paBXEoZIAhfZRJmo2jL",
#             "track_name": "Finalement",
#             "album_uri": "spotify:album:2KrRMJ9z7Xjoz1Az406UML",
#             "duration_ms": 166264,
#             "album_name": "Dancing Chords and Fireflies"
#         },
#     ],
# }

# })

for node in G.nodes(data=True):
    t = node[1]["node_type"]
    node_id(t, node[0])
    if t == "playlist":
        if node[1]["collaborative"] not in ("true", "false"):
            raise ValueError(f"collaborative is not a boolean: {node[1]['collaborative']}")
        node_features_by_type["playlist"] += [
            node[1]["num_followers"],
            node[1]["collaborative"] == 'true',
            node[1]["num_albums"],
            node[1]["num_tracks"],
            node[1]["num_edits"],
            node[1]["duration_ms"],
            node[1]["num_artists"]
        ]
    elif t == "track":
        distances = nx.single_source_shortest_path_length(G, node[0], cutoff=2)
        node_features_by_type["track"] += [[node[1]["duration"], len(distances)]]
    elif t == "artist":
        distances = nx.single_source_shortest_path_length(G, node[0], cutoff=2)
        node_features_by_type["artist"] += [[len(distances)]]
    elif t == "album":
        distances = nx.single_source_shortest_path_length(G, node[0], cutoff=2)
        node_features_by_type["album"] += [[len(distances)]]

edge_index_by_type = {
    ("track", "contains", "playlist"): [],
    ("track", "includes", "album"): [],
    ("track", "authors", "artist"): []
}

existing_edges = set()
for edge in G.edges(data=True):
    track_node = edge[0]
    other_node = edge[1]
    if "track" not in track_node:
        track_node, other_node = other_node, track_node

```

```

    if (track_node, other_node) in existing_edges:
        continue

    if G[edge[0]][edge[1]]["edge_type"] == "track-playlist":
        s_id = node_id("track", track_node)
        d_id = node_id("playlist", other_node)

        edge_index_by_type[("track", "contains", "playlist")] += [(s_id, d_id)]

    elif G[edge[0]][edge[1]]["edge_type"] == "track-album":
        s_id = node_id("track", track_node)
        d_id = node_id("album", other_node)

        edge_index_by_type[("track", "includes", "album")] += [(s_id, d_id)]

    elif G[edge[0]][edge[1]]["edge_type"] == "track-artist":
        s_id = node_id("track", track_node)
        d_id = node_id("artist", other_node)

        edge_index_by_type[("track", "authors", "artist")] += [(s_id, d_id)]

    existing_edges.add((track_node, other_node))

# construct HeteroData
hetero = torch_geometric.data.HeteroData()

# add initial node features
hetero["playlist"].x = torch.FloatTensor(node_features_by_type["playlist"].reshape(-1, len(node_features_by_type["playlist"])[0]))
hetero["track"].x = torch.FloatTensor(node_features_by_type["track"].reshape(-1, len(node_features_by_type["track"])[0]))
hetero["artist"].x = torch.FloatTensor(node_features_by_type["artist"].reshape(-1, len(node_features_by_type["artist"])[0]))
hetero["album"].x = torch.FloatTensor(node_features_by_type["album"].reshape(-1, len(node_features_by_type["album"])[0]))

# add edge indices
hetero["track", "contains", "playlist"].edge_index = torch.tensor(edge_index_by_type[("track", "contains", "playlist")]).t()
hetero["track", "includes", "album"].edge_index = torch.tensor(edge_index_by_type[("track", "includes", "album")]).t()
hetero["track", "authors", "artist"].edge_index = torch.tensor(edge_index_by_type[("track", "authors", "artist")]).t()

# post-processing
hetero = torch_geometric.transforms.ToUndirected()(hetero)
hetero = torch_geometric.transforms.NormalizeFeatures()(hetero)
return hetero

def ghetero2datasets(ghetero):
    """Split the dataset into train, validation and test sets."""
    transform = T.Compose([
        T.NormalizeFeatures(),
        T.RandomLinkSplit(
            num_val=0.1,
            num_test=0.1,
            disjoint_train_ratio=0.3,
            neg_sampling_ratio=2.0,
            add_negative_train_samples=False,
            edge_types=("track", "contains", "playlist"),
            rev_edge_types=("playlist", "rev_contains", "track"),
        )
    ])

    return transform(ghetero) # 3-tuple: data_train, data_val, data_test

```

## Preprocessing

We took a subset of our dataset by taking a subgraph of 5000 nodes.

```

def get_neigh_of_edge_type(G, edge_type, node):
    undirected_neigh = itertools.chain(G.neighbors(node), G.predecessors(node))
    return [
        n
        for n in undirected_neigh
        if (
            G.succ[node].get(n, dict()).get('edge_type', None) == edge_type or
            G.pred[node].get(n, dict()).get('edge_type', None) == edge_type
        )
    ]

def top_n_by_followers(G, n, node_type):
    """Get all nodes of type `node_type`."""
    playlists = [node for node in G.nodes(data=True) if node[1]["node_type"] == node_type]
    return sorted(playlists, key=lambda x: "num_followers" in x[1] and x[1]["num_followers"], reverse=True)[:n]

```

```

def get_smart_playlist_subset(G, playlists_to_keep):
    keep_nodes = set()
    for node in playlists_to_keep:
        keep_nodes.add(node[0])
        tracks = get_neigh_of_edge_type(G, "track-playlist", node[0])
        artists_and_albums = []

        for track in tracks:
            artists_and_albums += get_neigh_of_edge_type(G, "track-artist", track)
            artists_and_albums += get_neigh_of_edge_type(G, "track-album", track)

        keep_nodes = keep_nodes.union(set(tracks))
        keep_nodes = keep_nodes.union(set(artists_and_albums))
    return keep_nodes

def smart_split(G, splits=[100,500,1000,5000,10000]):
    ret = [None for _ in splits]
    for i in tqdm(splits):
        print(f"[{i}] started")
        start = time.time()
        playlists_to_keep = top_n_by_followers(G, i, "playlist")
        print(f"[{i}] got top n playlists in {time.time() - start} seconds")
        start = time.time()
        keep_nodes = get_smart_playlist_subset(G, playlists_to_keep)
        print(f"[{i}] finished getting neighbors in {time.time() - start} seconds")
        print(f"\t{len(keep_nodes)} nodes = {len(keep_nodes)/len(G.nodes)} % of graph")
        start = time.time()
        G_sub = nx.Graph(G.subgraph(keep_nodes))
        print(f"[{i}] finished subgraphing in {time.time() - start} seconds")
        start = time.time()
        splits[i] = G_sub
        print(f"[{i}] finished pickling in {time.time() - start} seconds")
    return ret

```

## Model

Check if cuda is available

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

GNN embedding prediction network (three SAGEConv layers)

```

class GNN(torch.nn.Module):
    def __init__(self, hidden_channels):
        super().__init__()
        self.conv1 = torch_geometric.nn.SAGEConv((-1, -1), hidden_channels, normalize=True, dropout=True, bias=True, dropout_prob=0.1)
        self.conv2 = torch_geometric.nn.SAGEConv((-1, -1), hidden_channels, normalize=True, dropout=True, bias=True, dropout_prob=0.1)
        self.conv3 = torch_geometric.nn.SAGEConv((-1, -1), hidden_channels)

        self.reset_parameters()

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index)
        x = torch.nn.functional.leaky_relu(x, negative_slope=0.1)
        x = self.conv2(x, edge_index)
        x = torch.nn.functional.leaky_relu(x, negative_slope=0.1)
        x = self.conv3(x, edge_index)
        return x

    def reset_parameters(self):
        self.conv1.reset_parameters()
        self.conv2.reset_parameters()
        self.conv3.reset_parameters()

```

Link predictor (predicts using dot product)

```

class LinkPredictor(torch.nn.Module):

    def __init__(self):
        super().__init__()

    def forward(self, x_track, x_playlist, track_playlist_edge):
        track_embedding = x_track[track_playlist_edge[0]]
        playlist_embedding = x_playlist[track_playlist_edge[1]]

        # Apply dot-product to get a prediction per supervision edge:
        return (track_embedding * playlist_embedding).sum(dim=-1)

```

## Full model

```

class HeteroModel(torch.nn.Module):
    def __init__(self, hidden_channels, node_features, metadata):
        super().__init__()
        self.node_lin = {
            k: torch.nn.Linear(v.shape[1], hidden_channels).to(device) for k, v in node_features.items()
        }
        self.gnn = GNN(hidden_channels).to(device)
        self.gnn = torch_geometric.nn.to_hetero(self.gnn, metadata=metadata).to(device)
        self.classifier = LinkPredictor().to(device)

    def embed(self, data):
        x_dict = {
            k: self.node_lin[k](v) for k, v in data.x_dict.items()
        }
        x_dict = self.gnn(x_dict, data.edge_index_dict)
        return x_dict

    def forward(self, data):
        x_dict = self.embed(data)
        pred = self.classifier(
            x_dict["track"],
            x_dict["playlist"],
            data["track", "contains", "playlist"].edge_label_index,
        )
        return pred

    def reset_parameters(self):
        for _, v in self.node_lin.items():
            torch.nn.init.xavier_uniform_(v.weight)
        self.gnn.reset_parameters()

def dummy_generator(source):
    for e in source:
        yield e

```

## Model train and test functions

```

outs = []

def test(model, data_test):
    with torch.no_grad():
        test_out = model(data_test.to(device)).to('cpu')
        truth = data_test["track", "contains", "playlist"].edge_label.to('cpu')

    test_loss = torch.nn.functional.mse_loss(
        test_out,
        truth
    )
    metric = BinaryAccuracy()
    metric.update(test_out, truth)
    return float(test_loss), metric.compute()

def train(model, train_loader, optimizer, batch_wrapper=dummy_generator):
    model.train()

    accuracy = 0

    total_examples = total_loss = 0
    for batch in batch_wrapper(train_loader):
        optimizer.zero_grad()

        out = model(batch)
        truth = batch["track", "contains", "playlist"].edge_label

        loss = torch.nn.functional.mse_loss(
            out, truth
        )
        loss.backward()
        optimizer.step()

        aute_gledam = out.to('cpu')

        outs.append(aute_gledam)

    metric = BinaryAccuracy()
    metric.update(aute_gledam, truth.to('cpu'))
    accuracy += metric.compute() * len(outs)

```

```

total_examples += len(out)
total_loss += float(loss) * len(out)

return total_loss / total_examples, accuracy / total_examples

```

## Train the model

```

G = load_graph()
ghetero = nx2hetero(G)
data_train, data_val, data_test = ghetero2datasets(ghetero)

# create training mask for playlist nodes
train_mask = torch.zeros(ghetero["playlist"].x.shape[0], dtype=torch.bool)
train_mask[torch.randperm(train_mask.shape[0])[:int(train_mask.shape[0]*0.8)]] = True

ghetero["playlist"].train_mask = train_mask

ghetero["playlist"].y = torch.LongTensor([1]*ghetero["playlist"].x.shape[0]).to(device)

model = HeteroModel(64, ghetero.x_dict, ghetero.metadata()).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.0001, weight_decay=0.00001)
edge_label_index = data_train["track", "contains", "playlist"].edge_label_index
edge_label = data_train["track", "contains", "playlist"].edge_label

train_loader = torch_geometric.loader.LinkNeighborLoader(
    data=data_train,
    num_neighbors=[-1],
    neg_sampling_ratio=0.5,
    edge_label_index=(("track", "contains", "playlist"), edge_label_index),
    edge_label=edge_label,
    batch_size=20000,
    shuffle=True,
    transform=T.ToDevice(device)
)

epoch = 2000
render_graph = True

losses = []
accuracies = []
test_losses = []
test_accuracies = []

epoch_iter = tqdm(range(epoch), unit='epoch', desc='Training', bar_format='{desc:<5.5}{percentage:3.0f}%|{bar:10}{r_bar}')
for i in epoch_iter:
    loss, accuracy = train(model, train_loader, optimizer)
    losses.append(loss)
    accuracies.append(accuracy)
    test_loss, test_acc = test(model, data_val)
    test_losses.append(test_loss)
    test_accuracies.append(test_acc)
    epoch_iter.set_postfix_str(f"Train Loss: {loss:.4f}, Train Accuracy {accuracy:.4f}, Valid Loss {test_loss:.4f}, Valid Accuracy {test_acc:.4f}")

```

## Render learning graph

```


plt.clf()
# add labels
plt.plot(np.arange(len(accuracies)), accuracies, label='Accuracy')
plt.plot(np.arange(len(losses)), losses, label='Loss')
plt.plot(np.arange(len(test_losses)), test_losses, label='Test Loss')
plt.plot(np.arange(len(test_accuracies)), test_accuracies, label='Test Accuracy')

# start plot at 0
plt.ylim(0, 1)
plt.legend()
plt.show()

```

[Colab paid products](#) - [Cancel contracts here](#)

---

 1m 45s    completed at 2:17 AM